


# Central Limit Theorem and Sampling Distributions: Gaussian Distributions

# Coin flip example

- We have a coin, and we flip it  $N$  times. We'll call it a sample. As we increase  $N$ , we get a better estimate of the true probability of heads.
- in our coins example, the true proportion, called  $p^*$  comes from all possible (infinite) coin flips. We never get to see this
- This of course depends on if our model describes the true generating process for the data, otherwise we can find a  $p^*$  given a population, but still have model mis-specification error
- if we are only given one (finite sample sized) replication, which is the situation in real life, we can only estimate a probability  $\hat{p}$
- In our idealized, simulated case we have many  $M$  replications, and thus samples, and we can now find the **distribution** of estimated probabilities  $\hat{p}$

# M replications of N coin tosses



# Sampling distribution

As we let  $M \rightarrow \infty$ , for any  $N$ , the distribution induced on  $\hat{p}$  is the empirical **sampling distribution of the estimator** of the true probability  $p^*$ .

We could use the sampling distribution to get confidence intervals on  $\hat{p}$ .


$\hat{p}$  itself is a mean value, a mean of the 1s corresponding to heads and 0s corresponding to the tails.

Thus the distribution of  $\hat{p}$  is the **sampling distribution of the mean** of each sample

# 200 replications of N coin tosses



- If the heads are given the value 1 and the tails are given the value 0 then the mean of the sample gives us the fraction of heads in each replication
- these means will be **different!** The fluctuations from one replication to another is called a **distribution**

200 reps, 491 flips per sample -->



# What if we increase the size of each sample?


1000 ( $\infty$ ) replications of N coin tosses, left: N=491, right: N=982



# Properties of this distribution



- We can average this fraction over the replications..a mean of sample means.
- We can also calculate the standard deviation of these means.
- How do these vary with the size of the sample?

# Distribution of Sample Means






# mean and standard deviation of sample means



# Sample means and the population mean

Let  $x_1, x_2, \dots, x_n$  be a sequence of IID values from random variable  $X$ , which has some distribution with finite mean  $\mu$ .

Define  $E_{\{R\}}(x)$  to be the average in the vertical direction (axis=0) of the value  $x$ , for example the average of all the first coin's 1s and 0s in the replications.



Then,

$$E_{\{R\}}(N \bar{x}) = E_{\{R\}}(x_1 + x_2 + \dots + x_N) = E_{\{R\}}(x_1) + E_{\{R\}}(x_2) + \dots + E_{\{R\}}(x_N)$$

In limit  $M \rightarrow \infty$  of replications, each and every one of the expectations in RHS can be replaced by the population mean  $\mu$  using the law of large numbers! Thus:

$$E_{\{R\}}(N \bar{x}) = N \mu$$

$$E_{\{R\}}(\bar{x}) = \mu$$

In limit  $M \rightarrow \infty$  of replications the expectation value of the sample means converges to the population mean.

Now let underlying distribution on the population have well defined mean  $\mu$  AND a well defined variance  $\sigma^2$ .

$$V_{\{R\}}(N \bar{x}) = V_{\{R\}}(x_1 + x_2 + \dots + x_N) = V_{\{R\}}(x_1) + V_{\{R\}}(x_2) + \dots + V_{\{R\}}(x_N)$$


Now in limit  $M \rightarrow \infty$ , each of the variances in the RHS can be replaced by the population variance using the law of large numbers! Thus:

$$V_{\{R\}}(N \bar{x}) = N \sigma^2$$

$$V(\bar{x}) = \frac{\sigma^2}{N}$$

# These sampling distributions are Gaussian

$$p(y|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\left(\frac{y-\mu}{2\sigma}\right)^2}$$




# The Central Limit Theorem (CLT)

Let  $x_1, x_2, \dots, x_n$  be a sequence of IID values from a random variable  $x$ . Suppose that  $x$  has the finite mean  $\mu$  AND finite variance  $\sigma^2$ . Then:

$$S_n = \frac{1}{n} \sum_{i=1}^n x_i, \text{ converges to}$$

$$S_n \sim N\left(\mu, \frac{\sigma^2}{n}\right) \text{ as } n \rightarrow \infty.$$

# $\log(\text{standard error})$ vs $\log(N)$



# What did we do here?

- we asked the question of a population: whats the mean of a bernoulli sequence of coins ( $H=1, T=0$ )
- we realized we only had samples to solve this on, so did it there
- then we asked, whats the distribution of these means
- it turns out to be gaussian
- with a mean equal to what the population mean would have been
- and a variance equal to the population mean divided by the sample size
- so that when the samples become the population their individual means converge to the population mean



# Meaning

- weight-watchers' study of 1000 people, average weight is 150 lbs with  $\sigma$  of 30lbs.
- Randomly choose many samples of 100 people each, the mean weights of those samples would cluster around 150lbs with a standard error of 3lbs.
- a different sample of 100 people with an average weight of 170lbs would be more than 6 standard errors beyond the population mean.

# Frequentist Statistics

Answers the question: **What is Data?** with

"data is a **sample** from an existing **population**"

- data is stochastic, variable, in the sense that you can draw different samples
- model the sample. The model may have parameters
- The parameters are considered **FIXED**, and there is a **true value** in our population
- However, we can only find parameters for our sample, since in real-life we usually only get to see one sample.
- If we could somehow access multiple samples, these parameters would vary from sample to sample

